

Edge Clouds Control Plane and Management Data Consistency Challenges: Position Paper for IEEE International Conference on Cloud Engineering, 2019

Bohdan Dobrelia
OpenStack platform
Red Hat
Poznan, Poland
bdobrelia@redhat.com

Abstract—Fog computing is emerging Cloud of (Edge) Clouds technology. Its control plane and deployments data synchronization is a major challenge. Autonomy requirements expect even the most distant edge sites always manageable, available for monitoring and alerting, scaling up/down, upgrading and applying security fixes. Whenever temporary disconnected sites are managed locally or centrally, some changes and data need to be eventually synchronized back to the central site(s) with having its merge-conflicts resolved for the central data hub(s). While some data needs to be pushed from the central site(s) to the Edge, which might require resolving data collisions at the remote sites as well. In this paper, we position the outstanding data synchronization problems for OpenStack platform becoming a cloud solution number one for fog computing. We define the inter-cloud operational invariants based on that Always Available autonomy requirement. We show that a causally consistent key value storage is the best match for the outlined operational invariants and there is a great opportunity for designing such a solution for Edge clouds. Finally, the paper brings the vision of unified tooling to solve the data synchronization problems the same way for infrastructure owners, IaaS cloud operators and tenants running workloads for PaaS, like OpenShift or Kubernetes deployed on top of Edge clouds.

Index Terms—Open source software, Edge computing, Distributed computing, System availability, Design

I. GLOSSARY

Aside of the established terms [3], we define a few more for the data processing and operational points of view:

A. Deployment Data

Data that represents the configuration of *cloudlets* [3], like API endpoints URI, or numbers of deployed *edge nodes* [3] in *edge clouds* [3]. That data represents the most recent state of a deployment.

B. Cloud Data

Represents the most recent ¹ internal and publicly visible state of cloudlets, like cloud users or virtual routers. Cloud

¹when there is unresolved data merging conflicts, the most recent state becomes the best known state

data also includes logs, performance and usage statistics, state of message queues and the contents of databases.

C. Control Plane

Corresponds to any operations performed via cloudlets API endpoints or command line tooling. For example, starting a virtual machine instance, or creating a cloud user. Such operations are typically initiated by cloud applications, tenants or operators.

D. Management Plane

Corresponds to administrative actions performed via configuration and lifecycle management systems. Such operations are typically targeted for cloudlets, like edge nodes, *edge data centers* [3], or edge clouds. E.g., upgrading or reconfiguring cloudlets in a *virtual data center* [3], or scaling up edge nodes. And typically initiated by cloud infrastructure owners. For some cases, like Baremetal-as-a-Service, tenants may as well initiate actions executed via the management plane. Collecting logs, performance and usage statistics for monitoring and alerting systems also represents the management plane operations, although it operates with the cloud data.

E. Always Available

The operational mode of the control and management planes that corresponds to the best for today choices for the *sticky available* [4] consistency models, e.g. *Real-Time Causal* [2], or *causal+* [1].

II. INTRODUCTION

III. ANALYSIS AND DISCUSSION

A. Autonomy Requirements

We define always available autonomy for cloudlets as the following strict requirements:

- any operations performed on cloudlets state ² fit data consistency models that allow the involved control/management planes operating as always available, and there is no read-only or blocking access limitations.
- cloudlets data can be modified at any given moment of time, despite of inter-cloudlets network connectivity ³.
- data can be synchronized eventually across cloudlets to/from ⁴ its adjacent aggregation edge layer, but not horizontally ⁵.
- data replication conflicts can be resolved automatically or by hand, and/or queued ⁶ for later processing ⁷.

B. Operational Invariants

To be always available, control and management planes of cloudlets should provide the following operational capabilities (*invariants* hereafter):

- Foo

C. Data Consistency Requirements

The operational invariants dictate inevitable presense of shared state ⁸ among cloudlets. We know [4] that *causally consistent* [1] [2] data accessing mechanisms is the best fit for the declared invariants of the always available control and management planes.

In OpenStack, there is yet such data storages supported as a data backend. is normally stored in databases, which known as

IV. CONCLUSION

REFERENCES

- [1] W. Lloyd, M. J. Freedman, M. Kaminsky, and D. G. Andersen, "Dont settle for eventual: Scalable causal consistency for wide-area storage with COPS," Proc. 23rd ACM Symposium on Operating Systems Principles (SOSP 11), Cascais, Portugal, October 2011.
- [2] P. Mahajan, L. Alvisi, and M. Dahlin. "Consistency, availability, and convergence," Technical Report TR-11-22, Univ. Texas at Austin, Dept. Comp. Sci., 2011.
- [3] The Linux Foundation, "Open Glossary of Edge Computing," [Online]. Available: <https://github.com/State-of-the-Edge/glossary>
- [4] K. Kingsbury, "Consistency Models," [Online]. Available: <https://jepson.io/consistency>

²despite the cloudlets aliveness or failure conditions

³for disconnected/partitioned cloudlets, data can be modified via local control/management plane, if exists and not failed. Despite the adjacent aggregation edge layer [3] global view and/or quorum requirements

⁴bi-directional data replication is not a strict requirement though. For some cases, it is acceptable to have state replicated only from aggregation edge layer to cloudlet(s) under its control/management, like virtual machine or hardware provisioning images data. Or otherwise, logs, performance and metering statistics can flow from cloudlets to aggregation edge layer

⁵also implies there is no horizontal data replication across neighbor aggregation edge layers. This corresponds to the acyclic graph (tree) topology

⁶depending on the numbers and allowed isolation periods of cloudlets under control/management, the disc and memory requirements for aggregation edge layers may vary drastically

⁷COPS [1] supports a similar failure mode for a datacenter: "the system administrator has two options: allow the queues to grow if the partition is likely to heal soon, or reconfigure COPS to no longer use the failed datacenter", except that the datacenter may be either failed or operating isolated long-time, and that we can have multiple such datacenters

⁸when we refer to just *data* or *state*, we do not differentiate either that is deployment or control data